

## Design Document

### Question 1 (file descriptors):

How are file descriptors implemented? What kernel data structures were created to manage file descriptors? Briefly describe the implementation of open, close, read, and write.

File descriptor is a structure containing a vnode pointer "v", an integer "flg", and an off\_t (64 bit integer) "os". v is the abstract representation of the file using by the file descriptor; flg indicates how to open the file, such as read-only which is represented by 0 and write-only which is represented by 1; os is the offset within the file that the file descriptor associated with.

A new field fdt (which is a fd\_table structure) is added to the thread structure. Each fd\_table contains a pointer to an array (array structure defined by kernel) of file descriptors called "fds" and an integer named "num\_fd". fds contains all the file descriptors of the thread, where their index in the array is their ID. Three file descriptor are created when a fdt is initialized. They are standard input, standard output, and standard error (fd 0, 1, 2). The I/O are operated through console device (con:). The fd\_num keeps track of the number of fd opened for the thread to prevent from open too many files.

### Open:

Open includes a char pointer "filename", an integer "flags", an integer mode, and an integer pointer "return\_val". After check for valid filename and flags, vfs\_open is called to open the file/device/other kernel object named by the pathname filename in mode specified by flags, and store contents into a vnode "v". If success, vfs\_open will return 0. fd\_create will then be called with vnode v, flag flags, and an offset of 0, which will create a file descriptor "fd". fd will then be added to current thread's fdt, appending to the end of the fds array. This is done by calling fd\_table\_add\_fd, which returns the added fd's ID upon success and returns -1 when failed. If added successfully, the fd ID will be stored at the address pointed to by return\_val, and 0 will be returned to indicate a successful Open of a file. Current thread's fd\_num field will increase by 1, to avoid open more file than the limit number of files sys161 can open. Whenever an error happens, the open function will return -1, and the err number indicating the type of error is stored in the address pointed by return\_val.

### Close:

Close includes an integer "fd" and an integer pointer "return\_val". It first checks for valid input of fd, such as deadbeef or a fd ID bigger than the length of array fds. Then fd\_table\_rm\_fd will be called, which "removes" the fd from

current thread's fdt's fds by calling `vfs_close` and setting `fd` to null in the array. The fdt's `num_fd` is decreased by 1 every time a `fd` is removed. If success, returns 0; if not, returns 1, and the close function will return -1 and store the error number in the address pointed by `return_val`. If the close function successfully close the file descriptor, it will return 0.

#### Read:

Read takes in an integer "`fd`", a void pointer "`buf`", a `size_t` `buflen`, and an integer pointer "`return_val`". It checks for valid address, valid `fd`, and make sure the function is not write-only. An `uio` structure is initialized to store data read by `VOP_READ` from the given `fd`'s `vnode`. If any of the above condition failed, function will return -1 and and store the error number in the address pointed to by `return_val`. `fd`'s offset will be set to `u`'s offset to indicate the number of byte read already. If `u`'s `uio_resid` equals 0, no more bytes are needed to read, and the number of bytes read already (which equals to `buflen`) will be stored in the address `return_val` points to. If there is still bytes left to be read, amount of bytes left to be read (`buflen-u.uio_resid`) will be stored. Upon success, Read function returns 0.

#### Write:

Write takes in an integer "`fd`", a const void pointer "`buf`", a `size_t` `nbytes`, and an integer pointer "`return_val`". It checks for valid address, valid `fd`, and make sure the function is not read-only. An `uio` structure is initialized to store data to be write by `VOP_WRITE` to the given `fd`'s `vnode`. If any of the above condition failed, function will return -1 and and store the error number in the address pointed to by `return_val`. `fd`'s offset will be set to `u`'s offset to indicate the number of byte wrote already. If `u`'s `uio_resid` equals 0, no more bytes are required to write, and the number of bytes wrote already (which equals to `nbytes`) will be stored in the address `return_val` points to. If there is still bytes left to be write, the amount (`nbytes-u.uio_resid`) will be stored. Upon success, Write function returns 0.